

Problem set 4

2024-10-04

In the next problem set, we plan to explore the relationship between COVID-19 death rates and vaccination rates across US states by visually examining their correlation. This analysis will involve gathering COVID-19 related data from the CDC's API and then extensively processing it to merge the various datasets. Since the population sizes of states vary significantly, we will focus on comparing rates rather than absolute numbers. To facilitate this, we will also source population data from the US Census to accurately calculate these rates.

In this problem set we will learn how to extract and wrangle data from the data US Census and CDC APIs.

1. Get an API key from the US Census at https://api.census.gov/data/key_signup.html. You can't share this public key. But your code has to run on a TFs computer. Assume the TF will have a file in their working directory named `census-key.R` with the following one line of code:

```
census_key <- "A_CENSUS_KEY_THAT_WORKS"
```

Write a first line of code for your problem set that defines `census_key` by running the code in the file `census-key.R`.

```
source('census-key.R')
```

2. The [US Census API User Guide](#) provides details on how to leverage this valuable resource. We are interested in vintage population estimates for years 2020 and 2021. From the documentation we find that the *endpoint* is:

```
url <- "https://api.census.gov/data/2021/pep/population"
```

Use the `httr2` package to construct the following GET request.

```
https://api.census.gov/data/2021/pep/population?get=POP_2020,POP_2021,NAME&for=state:*&key=Y
```

Create an object called `request` of class `httr2_request` with this URL as an endpoint. Hint: Print out `request` to check that the URL matches what we want.

```
library(httr2)
request <- request(url) |> req_url_query(get = I("POP_2020,POP_2021,NAME"),
                                         `for` = I("state:*"),
                                         key = census_key)
```

3. Make a request to the US Census API using the `request` object. Save the response to an object named `response`. Check the response status of your request and make sure it was successful. You can learn about *status codes* [here](#).

```
response <- request |> req_perform()
```

4. Use a function from the `httr2` package to determine the content type of your response.

```
library('httr2')
request |> req_perform() |> resp_content_type()
```

```
[1] "application/json"
```

5. Use just one line of code and one function to extract the data into a matrix. Hints: 1) Use the `resp_body_json` function. 2) The first row of the matrix will be the variable names and this OK as we will fix in the next exercise.

```
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

```
filter, lag
```

The following objects are masked from 'package:base':

```
intersect, setdiff, setequal, union
```

```
library(janitor)
```

Attaching package: 'janitor'

The following objects are masked from 'package:stats':

```
chisq.test, fisher.test
```

```
library(tidyr)
library(stringr)
population <- response |> resp_body_json(simplifyVector = TRUE)
```

6. Examine the `population` matrix you just created. Notice that 1) it is not tidy, 2) the column types are not what we want, and 3) the first row is a header. Convert `population` to a tidy dataset. Remove the state ID column and change the name of the column with state names to `state_name`. Add a column with state abbreviations called `state`. Make sure you assign the abbreviations for DC and PR correctly. Hint: Use the **janitor** package to make the first row the header.

```
library(tidyverse)
library(janitor)

population <- population |>
  row_to_names(1)|>
  as_tibble()|>
  select(-state)|>
  rename(state_name=NAME) |>
  pivot_longer(-state_name, names_to = 'year', values_to = 'population')|>
  mutate(year = str_remove(year, "POP_"))|>
  mutate(across(-state_name, as.numeric))|>
  mutate(state= state.abb[match(state_name, state.name)]) |>
  mutate(state = case_when (state_name == "Puerto Rico" ~ "PR",
                           state_name == "District of Columbia" ~ "DC",
                           .default = state))

head(population)
```

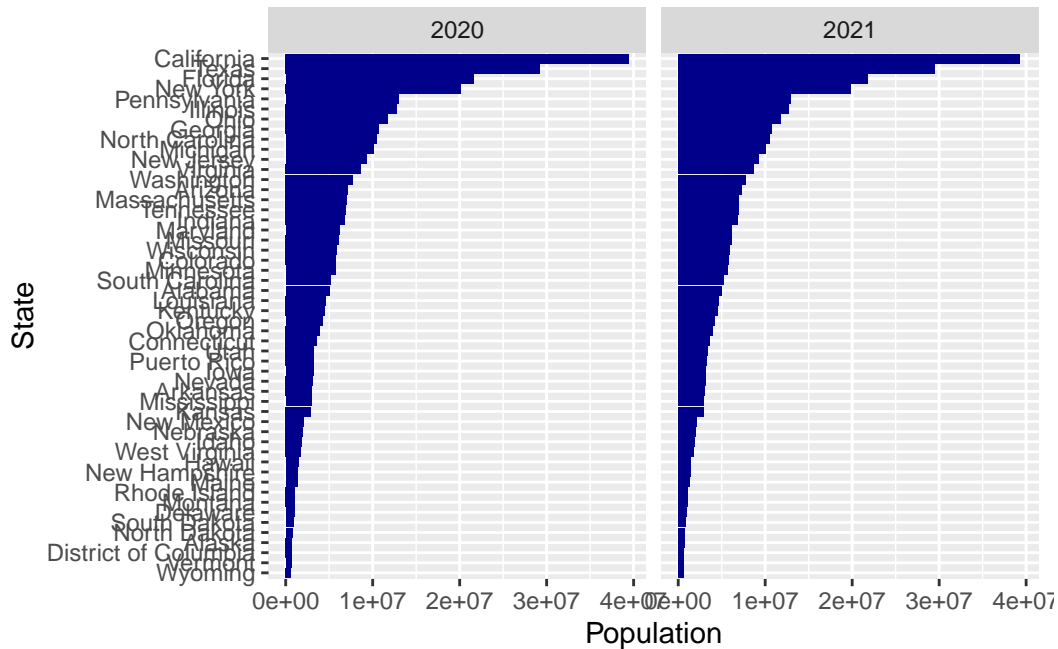
```
# A tibble: 6 x 4
  state_name  year population state
  <chr>      <dbl>      <dbl> <chr>
```

1	Oklahoma	2020	3962031	OK
2	Oklahoma	2021	3986639	OK
3	Nebraska	2020	1961455	NE
4	Nebraska	2021	1963692	NE
5	Hawaii	2020	1451911	HI
6	Hawaii	2021	1441553	HI

```
#population <- population |> ## Use janitor row to names function
# convert to tibble
# remove stat column
# rename state column to state_name
# use pivot_longer to tidy
# remove POP_ from year
# parse all relevant columns to numeric
# add state abbreviations using state.abb variable
# use case_when to add abbreviations for DC and PR
```

7. As a check, make a barplot of states' 2021 and 2022 populations. Show the state names in the y-axis ordered by population size. Hint: You will need to use `reorder` and use `facet_wrap`.

```
population |>
  ggplot(aes(x = reorder(state_name, population), y = population))+
  geom_col(fill='darkblue') +
  coord_flip()+
  facet_wrap(~year)+
  labs(x = "State", y = "Population")
```



```
# population |>
# reorder state
# assign aesthetic mapping
# use geom_col to plot barplot
# flip coordinates
# facet by year
```

8. The following URL:

```
url <- "https://github.com/datasciencelabs/2024/raw/refs/heads/main/data/regions.json"
```

points to a JSON file that lists the states in the 10 Public Health Service (PHS) defined by CDC. We want to add these regions to the `population` dataset. To facilitate this create a data frame called `regions` that has two columns `state_name`, `region`, `region_name`. One of the regions has a long name. Change it to something shorter.

```
library(jsonlite)
library(purrr)
regions <- fromJSON(url,simplifyDataFrame = FALSE)
regions <- map_df(regions, function(x)
  data.frame(region = x$region, region_name = x$region_name, state_name = x$states)) |>
  mutate(region_name = ifelse(region_name == "New York and New Jersey, Puerto Rico, Virgin I",
  head(regions)
```

	region	region_name	state_name
1	1	New England	Connecticut
2	1	New England	Maine
3	1	New England	Massachusetts
4	1	New England	New Hampshire
5	1	New England	Rhode Island
6	1	New England	Vermont

```
# regions <- use jsonlit JSON parser
# regions <- convert list to data frame. You can use map_df in purrr package
```

9. Add a region and region name columns to the population data frame.

```
population <- left_join(population, regions, by = "state_name")
head(population)
```

```
# A tibble: 6 x 6
  state_name year population state region region_name
  <chr>      <dbl>      <dbl> <chr>   <int> <chr>
1 Oklahoma  2020      3962031 OK       6 South Central
2 Oklahoma  2021      3986639 OK       6 South Central
3 Nebraska  2020      1961455 NE       7 Central Plains
4 Nebraska  2021      1963692 NE       7 Central Plains
5 Hawaii    2020      1451911 HI       9 Pacific
6 Hawaii    2021      1441553 HI       9 Pacific
```

10. From reading <https://data.cdc.gov/> we learn the endpoint <https://data.cdc.gov/resource/pwn4-m3yp> provides state level data from SARS-COV2 cases. Use the **httr2** tools you have learned to download this into a data frame. Is all the data there? If not, comment on why.

```
api <- "https://data.cdc.gov/resource/pwn4-m3yp.json"
cases_raw <- request(api) |> req_perform() |> resp_body_json(simplifyVector = TRUE)
head(cases_raw)
```

	date_updated	state	start_date	end_date
1	2023-02-23T00:00:00.000	AZ	2023-02-16T00:00:00.000	2023-02-22T00:00:00.000
2	2022-12-22T00:00:00.000	LA	2022-12-15T00:00:00.000	2022-12-21T00:00:00.000
3	2023-02-23T00:00:00.000	GA	2023-02-16T00:00:00.000	2023-02-22T00:00:00.000
4	2023-03-30T00:00:00.000	LA	2023-03-23T00:00:00.000	2023-03-29T00:00:00.000
5	2023-02-02T00:00:00.000	LA	2023-01-26T00:00:00.000	2023-02-01T00:00:00.000
6	2023-03-23T00:00:00.000	LA	2023-03-16T00:00:00.000	2023-03-22T00:00:00.000

	tot_cases	new_cases	tot_deaths	new_deaths	new_historic_cases
1	2434631.0	3716.0	33042.0	39.0	23150
2	1507707.0	4041.0	18345.0	21.0	21397
3	3061141.0	5298.0	42324.0	88.0	6800
4	1588259.0	2203.0	18858.0	23.0	5347
5	1548508.0	5725.0	18572.0	47.0	4507
6	1580709.0	1961.0	18835.0	35.0	2239

	new_historic_deaths
1	0
2	0
3	0
4	0
5	0
6	0

No, data is not all there. We see exactly 1,000 rows. We should be seeing over 52×3 rows per state.

- The reason you see exactly 1,000 rows is because CDC has a default limit. You can change this limit by adding `$limit=10000000000` to the request. Rewrite the previous request to ensure that you receive all the data. Then wrangle the resulting data frame to produce a data frame with columns `state`, `date` (should be the end date) and `cases`. Make sure the cases are numeric and the dates are in `Date` ISO-8601 format.

```
api <- "https://data.cdc.gov/resource/pwn4-m3yp.json?$limit=10000000000"
cases_raw <- request(api) |> req_perform() |> resp_body_json(simplifyVector = TRUE)
cases <- cases_raw|>
  mutate(date = as.Date(end_date), cases = as.numeric(new_cases))|>
  select(state, date, cases )
```

- For 2020 and 2021, make a time series plot of cases per 100,000 versus time for each state. Stratify the plot by region name. Make sure to label you graph appropriately.

```
cases|> filter(year(date) %in% c(2020, 2021))|>
  mutate(year = year(date))|>
  left_join(population, by =c("state",'year')) |>
  mutate(cases_per_100k = (cases / population) * 100000) |>
  filter(!is.na(cases_per_100k))|>
  ggplot(aes(x = date, y = cases_per_100k, color = state_name))+
  geom_line()+
  facet_wrap(~region_name) +
  theme(axis.text.y = element_text(angle = 0, hjust = 3, vjust = 2)) +
```

```

labs(x = "Date",
     y = "Cases per 100,000") + # Labels
theme_minimal() + # Use a minimal theme
theme(
  panel.spacing = unit(0.1, "lines"),
  legend.position = "bottom",
  axis.text.x = element_text(size = 8, angle = 45, hjust = 1),
  axis.text.y = element_text(size = 8),
  legend.key.size = unit(0.2, "cm"),
  legend.text = element_text(size = 7)
)

```

